

In [1]: *## Preparation*

```
# Import the necessary Libraries
import torch
import torch.nn as nn
from sklearn.model_selection import train_test_split
from torchvision import models, datasets, transforms
from torch.utils.data import DataLoader, Subset
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay, classification_report
import matplotlib.pyplot as plt
import numpy as np
import os
import optuna
import time

## Global Var

dataset_path = os.getcwd() + '/dataset/classified'

class_labels = ["Normal", "Cataract"]
```

In [2]: `print(f"Is using CUDA? {torch.cuda.is_available()}")` *# Should return True if CUDA*  
`print(torch.version.cuda)` *# Check the CUDA version PyTorch is using*  
`print(torch.cuda.current_device())` *# Check CUDA device used*

```
Is using CUDA? True
12.6
0
```

In [3]: *# Augmentation*

```
transform = transforms.Compose([
    transforms.Resize((224, 224)),
    transforms.RandomRotation(15),
    transforms.RandomHorizontalFlip(p=0.5),
    transforms.RandomCrop(224, padding=4),
    transforms.ToTensor(),
    transforms.Normalize(mean=[0.5, 0.5, 0.5], std=[0.5, 0.5, 0.5]) # Normalize
])

# Load Dataset
ds = datasets.ImageFolder(root=dataset_path, transform=transform)

indices = list(range(len(ds)))
# Labels = [ds.targets[i] for i in indices]

# Split into train and test dataset
train_indices, test_indices = train_test_split(indices, test_size=0.2, random_state=42)

train_ds = Subset(ds, train_indices)
test_ds = Subset(ds, test_indices)

train_loader = DataLoader(train_ds, batch_size=64, shuffle=True, pin_memory=True, num_workers=4)
test_loader = DataLoader(test_ds, batch_size=64, shuffle=True, pin_memory=True, num_workers=4)
```

```

total_samples = len(train_ds) + len(test_ds)

print(f"Train size: {(len(train_ds) / total_samples) * 100:.2f}%, Test size: {(len(
print(f"Total samples: {total_samples}, Train size: {len(train_ds)}, Test size: {le

```

Train size: 79.98%, Test size: 20.02%

Total samples: 1159, Train size: 927, Test size: 232

```

In [4]: def denormalize(tensor, mean=None, std=None):
        if std is None:
            std = [0.5, 0.5, 0.5]
        if mean is None:
            mean = [0.5, 0.5, 0.5]
        mean = torch.tensor(mean).view(3, 1, 1)
        std = torch.tensor(std).view(3, 1, 1)
        return tensor * std + mean # Reverse normalization

# Get a batch of images
dataiter = iter(train_loader)
images, labels = next(dataiter)

# Select one image
img = images[0]
label = labels[0].item()


# Denormalize image
img = denormalize(img)

# Convert from Tensor (C, H, W) to NumPy (H, W, C)
img = np.transpose(img.numpy(), (1, 2, 0))

# Plot the image
plt.imshow(img)
plt.title(f"Label: {class_labels[label]}") # Display label
plt.axis("off")
plt.show()

```

Label: Normal



```
In [5]: # Optuna Hyperparameter

best_trial = None
best_model = None

torch.backends.cudnn.benchmark = True

device_name = "cuda" if torch.cuda.is_available() else "cpu"
use_amp = device_name == "cuda"

def objective(trial: optuna.Trial) -> float:
    device = torch.device(device_name)

    # Define hyperparameters
    lr = trial.suggest_float("lr", 1e-3, 1e-1, log=True)
    dropout_rate = trial.suggest_float("dropout", 0.2, 0.5)
    optimizer_name = trial.suggest_categorical("optimizer", ["AdamW", "SGD"])
    num_epochs = trial.suggest_int("num_epochs", 5, 10) # Fixed at 5 epochs

    model = models.efficientnet_b0(progress=True, weights=models.EfficientNet_B0_We

    for param in model.features[:-2].parameters():
        param.requires_grad = False

    number_of_features = model.classifier[1].in_features
    model.classifier = nn.Sequential(
        nn.Linear(number_of_features, 128),
        nn.ReLU(),
        nn.Dropout(dropout_rate),
        nn.Linear(128, 1)
```

```

).to(device)

model.to(device)

# Define Loss and optimizer
criterion = nn.BCEWithLogitsLoss()

optimizer = torch.optim.AdamW(model.parameters(), lr=lr, weight_decay=1e-4)
if optimizer_name == "SGD":
    optimizer = torch.optim.SGD(model.parameters(), lr=lr, momentum=0.9)

scheduler = torch.optim.lr_scheduler.CosineAnnealingLR(optimizer, T_max=10)

train_losses = []
train_accuracies = []

scaler = torch.amp.GradScaler(device=device_name)

# Training Loop
for epoch in range(num_epochs):

    model.train()
    total_loss, correct, total = 0, 0, 0

    for images, labels in train_loader:
        images, labels = images.to(device), labels.float().to(device).unsqueeze(0)

        optimizer.zero_grad()

        with torch.amp.autocast("cuda", enabled=use_amp):
            outputs = model(images)
            loss = criterion(outputs, labels)

        scaler.scale(loss).backward()
        scaler.step(optimizer)
        scaler.update()

        total_loss += loss.item()
        preds = (torch.sigmoid(outputs) > 0.5).float()
        correct += (preds == labels).sum().item()
        total += labels.size(0)

    scheduler.step()

    epoch_loss = total_loss / len(train_loader)
    epoch_acc = correct / total * 100

    train_losses.append(epoch_loss)
    train_accuracies.append(epoch_acc)

    print(f"Epoch [{epoch+1}/{num_epochs}], Loss: {epoch_loss:.4f}, Accuracy: {epoch_acc:.4f}")

    trial.set_user_attr("train_losses", train_losses)
    trial.set_user_attr("train_accuracies", train_accuracies)

# Pruning: Stop bad trials early

```

```

        trial.report(epoch_acc, epoch)
        if trial.should_prune():
            raise optuna.exceptions.TrialPruned()

    # Evaluate Model
    model.eval()
    correct, total = 0, 0
    trial_preds = []
    trial_labels = []

    with torch.no_grad():
        for images, labels in test_loader:
            images, labels = images.to(device), labels.float().to(device).unsqueeze(1)
            outputs = model(images)
            preds = (torch.sigmoid(outputs) > 0.5).float()

            correct += (preds == labels).sum().item()
            total += labels.size(0)

            trial_preds.extend(preds.cpu().numpy())
            trial_labels.extend(labels.cpu().numpy())

    test_acc = correct / total * 100
    print(f"🎯 Test Accuracy: {test_acc:.2f}%")

    # **Save best model globally**
    global best_model, best_trial

    try:
        best_trial = trial.study.best_trial
        best_value = trial.study.best_value
    except ValueError:
        best_value = float('-inf')

    if best_model is None or (best_trial is not None and test_acc > best_value):
        best_model = model.state_dict()

    best_acc = trial.study.user_attrs.get("best_accuracy", 0)

    if test_acc > best_acc:
        study.set_user_attr("best_accuracy", test_acc)

        trial_preds = np.array(trial_preds).flatten().tolist()
        trial_labels = np.array(trial_labels).flatten().tolist()

        study.set_user_attr("best_preds", trial_preds)
        study.set_user_attr("best_labels", trial_labels)

    return test_acc

```

```

In [6]: study = optuna.create_study(
        direction="maximize",
        study_name=f"hyperparam cataract classifier_{int(time.time())}",
        pruner=optuna.pruners.MedianPruner(),
        storage="sqlite:///optuna.db",
        load_if_exists=True
    )

```

```
)

study.optimize(
    objective,
    n_trials=20,
    n_jobs=4,
    show_progress_bar=True
)

acc = study.best_value

print("\nBest Hyperparameters:", study.best_params)
print("\nAccuracy:", acc)
```

```
[I 2025-03-13 02:15:18,379] A new study created in RDB with name: hyperparam cataract classifier_1741806917
0%|          | 0/20 [00:00<?, ?it/s]
```

Epoch [1/10], Loss: 0.3981, Accuracy: 83.50%  
Epoch [1/6], Loss: 1.1209, Accuracy: 70.77%  
Epoch [1/10], Loss: 0.6950, Accuracy: 49.08%  
Epoch [1/10], Loss: 0.3609, Accuracy: 81.77%  
Epoch [2/10], Loss: 0.2310, Accuracy: 91.69%  
Epoch [2/6], Loss: 0.3079, Accuracy: 87.70%  
Epoch [2/10], Loss: 0.6490, Accuracy: 70.33%  
Epoch [2/10], Loss: 0.2020, Accuracy: 91.91%  
Epoch [3/6], Loss: 0.2354, Accuracy: 90.83%  
Epoch [3/10], Loss: 0.2082, Accuracy: 93.31%  
Epoch [3/10], Loss: 0.6113, Accuracy: 70.66%  
Epoch [3/10], Loss: 0.1619, Accuracy: 94.93%  
Epoch [4/6], Loss: 0.1987, Accuracy: 92.77%  
Epoch [4/10], Loss: 0.1583, Accuracy: 94.28%  
Epoch [4/10], Loss: 0.5920, Accuracy: 70.66%  
Epoch [4/10], Loss: 0.1492, Accuracy: 94.17%  
Epoch [5/10], Loss: 0.1396, Accuracy: 94.17%  
Epoch [5/10], Loss: 0.5830, Accuracy: 70.66%  
Epoch [5/6], Loss: 0.2031, Accuracy: 94.17%  
Epoch [5/10], Loss: 0.1069, Accuracy: 96.33%  
Epoch [6/10], Loss: 0.1233, Accuracy: 95.15%  
Epoch [6/10], Loss: 0.5658, Accuracy: 70.66%  
Epoch [6/6], Loss: 0.1655, Accuracy: 93.42%  
Epoch [6/10], Loss: 0.0855, Accuracy: 97.09%

🔴 Test Accuracy: 93.10%

[I 2025-03-13 02:19:41,030] Trial 0 finished with value: 93.10344827586206 and parameters: {'lr': 0.02989743682562383, 'dropout': 0.41864583212524387, 'optimizer': 'AdamW', 'num\_epochs': 6}. Best is trial 0 with value: 93.10344827586206.

Epoch [7/10], Loss: 0.0914, Accuracy: 96.87%  
Epoch [7/10], Loss: 0.5619, Accuracy: 70.66%  
Epoch [7/10], Loss: 0.0821, Accuracy: 96.87%  
Epoch [1/8], Loss: 0.3949, Accuracy: 81.12%  
Epoch [8/10], Loss: 0.0684, Accuracy: 97.30%  
Epoch [8/10], Loss: 0.5485, Accuracy: 70.66%  
Epoch [8/10], Loss: 0.0629, Accuracy: 97.73%  
Epoch [2/8], Loss: 0.2255, Accuracy: 90.18%  
Epoch [9/10], Loss: 0.5536, Accuracy: 70.66%  
Epoch [9/10], Loss: 0.0575, Accuracy: 97.63%  
Epoch [9/10], Loss: 0.0544, Accuracy: 98.17%  
Epoch [10/10], Loss: 0.5460, Accuracy: 70.66%  
Epoch [3/8], Loss: 0.1722, Accuracy: 93.64%  
Epoch [10/10], Loss: 0.0779, Accuracy: 96.98%  
Epoch [10/10], Loss: 0.0515, Accuracy: 98.60%

🔴 Test Accuracy: 75.00%

[I 2025-03-13 02:21:46,974] Trial 3 finished with value: 75.0 and parameters: {'lr': 0.0020507919934963925, 'dropout': 0.41314315121606576, 'optimizer': 'SGD', 'num\_epochs': 10}. Best is trial 0 with value: 93.10344827586206.

🔴 Test Accuracy: 94.40%

[I 2025-03-13 02:21:47,898] Trial 1 finished with value: 94.39655172413794 and parameters: {'lr': 0.008650262291857513, 'dropout': 0.3336420482173461, 'optimizer': 'AdamW', 'num\_epochs': 10}. Best is trial 1 with value: 94.39655172413794.

Epoch [4/8], Loss: 0.1351, Accuracy: 94.71%

🔴 Test Accuracy: 92.24%

[I 2025-03-13 02:21:58,049] Trial 2 finished with value: 92.24137931034483 and parameters: {'lr': 0.003966599858683413, 'dropout': 0.39308162601981733, 'optimizer': 'AdamW', 'num\_epochs': 10}. Best is trial 1 with value: 94.39655172413794.

Epoch [1/7], Loss: 0.5831, Accuracy: 68.93%  
Epoch [1/6], Loss: 0.6680, Accuracy: 63.86%  
Epoch [5/8], Loss: 0.1137, Accuracy: 95.36%  
Epoch [1/5], Loss: 3.9171, Accuracy: 55.34%  
Epoch [2/7], Loss: 0.3753, Accuracy: 82.74%  
Epoch [2/6], Loss: 0.5928, Accuracy: 70.66%  
Epoch [6/8], Loss: 0.0817, Accuracy: 97.52%  
Epoch [2/5], Loss: 0.5025, Accuracy: 76.27%  
Epoch [3/7], Loss: 0.2303, Accuracy: 91.26%  
Epoch [3/6], Loss: 0.5432, Accuracy: 70.66%  
Epoch [7/8], Loss: 0.0686, Accuracy: 97.52%  
Epoch [3/5], Loss: 0.2874, Accuracy: 89.32%  
Epoch [4/7], Loss: 0.1878, Accuracy: 92.66%  
Epoch [4/6], Loss: 0.5173, Accuracy: 70.66%  
Epoch [8/8], Loss: 0.0629, Accuracy: 97.30%  
Epoch [4/5], Loss: 0.2010, Accuracy: 92.13%

🌀 Test Accuracy: 92.67%

[I 2025-03-13 02:24:24,624] Trial 4 finished with value: 92.67241379310344 and parameters: {'lr': 0.0037016668606250436, 'dropout': 0.4264276214041657, 'optimizer': 'AdamW', 'num\_epochs': 8}. Best is trial 1 with value: 94.39655172413794.

Epoch [5/7], Loss: 0.1427, Accuracy: 95.04%  
Epoch [5/6], Loss: 0.4955, Accuracy: 71.31%  
[I 2025-03-13 02:24:27,893] Trial 6 pruned.  
Epoch [5/5], Loss: 0.2025, Accuracy: 92.45%  
[I 2025-03-13 02:24:48,061] Trial 7 pruned.  
Epoch [1/9], Loss: 3.5660, Accuracy: 61.92%  
[I 2025-03-13 02:24:58,281] Trial 8 pruned.  
Epoch [6/7], Loss: 0.1168, Accuracy: 95.25%  
Epoch [1/10], Loss: 0.5652, Accuracy: 70.44%  
[I 2025-03-13 02:25:00,242] Trial 9 pruned.  
Epoch [1/6], Loss: 0.6901, Accuracy: 53.51%  
[I 2025-03-13 02:25:23,363] Trial 10 pruned.  
Epoch [1/7], Loss: 0.4230, Accuracy: 78.53%  
[I 2025-03-13 02:25:30,475] Trial 11 pruned.  
Epoch [7/7], Loss: 0.1106, Accuracy: 95.47%  
[I 2025-03-13 02:25:31,643] Trial 5 pruned.  
Epoch [1/9], Loss: 0.6835, Accuracy: 55.77%  
[I 2025-03-13 02:25:33,617] Trial 12 pruned.  
Epoch [1/8], Loss: 0.4550, Accuracy: 79.61%  
[I 2025-03-13 02:25:58,307] Trial 13 pruned.  
Epoch [1/8], Loss: 0.7141, Accuracy: 72.60%  
[I 2025-03-13 02:26:02,427] Trial 14 pruned.  
Epoch [1/8], Loss: 0.5978, Accuracy: 78.10%  
[I 2025-03-13 02:26:03,555] Trial 15 pruned.  
Epoch [1/5], Loss: 0.6299, Accuracy: 71.84%  
[I 2025-03-13 02:26:07,526] Trial 16 pruned.  
Epoch [1/5], Loss: 0.6042, Accuracy: 73.46%  
[I 2025-03-13 02:26:30,627] Trial 17 pruned.  
Epoch [1/5], Loss: 0.9715, Accuracy: 72.92%  
[I 2025-03-13 02:26:33,234] Trial 18 pruned.  
Epoch [1/5], Loss: 1.3657, Accuracy: 64.29%  
[I 2025-03-13 02:26:33,883] Trial 19 pruned.

Best Hyperparameters: {'lr': 0.008650262291857513, 'dropout': 0.3336420482173461, 'optimizer': 'AdamW', 'num\_epochs': 10}



Accuracy: 94.39655172413794


```
In [7]: best_trial = study.best_trial
best_train_losses = best_trial.user_attrs.get("train_losses", [])
best_train_accuracies = best_trial.user_attrs.get("train_accuracies", [])

plt.figure(figsize=(10, 5))

# Loss Graph
plt.subplot(1, 2, 1)
plt.plot(best_train_losses, label="Loss", marker="o", linestyle="-", color="b")
plt.xlabel("Epoch")
plt.ylabel("Loss")
plt.title("Training Loss per Epoch")
plt.legend()

# Accuracy Graph
plt.subplot(1, 2, 2)
plt.plot(best_train_accuracies, label="Accuracy (%)", marker="o", linestyle="-", co
plt.xlabel("Epoch")
plt.ylabel("Accuracy")
plt.title("Training Accuracy per Epoch")
plt.legend()

plt.show()
```




```
In [8]: best_preds = np.array(study.user_attrs.get("best_preds", []))
best_labels = np.array(study.user_attrs.get("best_labels", []))

conf_matrix = confusion_matrix(best_preds, best_labels)
ConfusionMatrixDisplay(conf_matrix, display_labels=class_labels).plot(cmap=plt.cm.B
cr = classification_report(best_preds, best_labels)
```

```
print(f"Accuracy: {acc:.2f}%")
print(cr)
```

Accuracy: 94.40%

	precision	recall	f1-score	support
0.0	0.95	0.98	0.96	169
1.0	0.93	0.86	0.89	63
accuracy			0.94	232
macro avg	0.94	0.92	0.93	232
weighted avg	0.94	0.94	0.94	232



```
In [9]: ## save
output_model_path = f"output/checkpoint-{acc}-hyperparam.pth"

torch.save({
    "model_state_dict": study.user_attrs.get("best_model_state"), # Best model wei
    "optimizer_state_dict": study.user_attrs.get("best_optimizer_state"), # Best o
    "best_hyperparameters": study.best_params,
    "best_accuracy": study.best_value,
    "best_train_losses": study.user_attrs.get("best_train_losses", []), # Best tra
    "best_train_accuracies": study.user_attrs.get("best_train_accuracies", []), #
}, output_model_path)
```

```
In [10]: # checkpoint = torch.load("checkpoint.pth")
# model.load_state_dict(checkpoint['model_state_dict'])
# optimizer.load_state_dict(checkpoint['optimizer_state_dict'])
# model.eval()
```

